
Efficient Estimation of Dynamic Conditional Correlation Models

Euan Dowers

June 28, 2017

Contents

1	Introduction	3
2	Dynamic Conditional Correlation Model	3
3	Shrinkage Estimation for Large Unconditional Correlation Matrices	4
4	DCC Model Estimation	6
4.1	Cholesky Decomposition	6
4.2	Updating the Cholesky Decomposition	7
5	Monte Carlo Simulation	7
5.1	Setup of Simulations	7
5.2	Results	8
6	Conclusion	9

1 Introduction

Estimating the correlation structure of a large number of financial assets is an important task for mean-variance optimisation, but is also an extremely difficult task, especially when the requirement is added that this correlation structure varies over time.

The aim of this paper is to efficiently estimate the Dynamic Conditional Correlation model in large dimensions (both in time and the number of assets included in the model), using two innovations. The first, due to Engle, Ledoit and Wolf [?], is to use the linear shrinkage method of Ledoit and Wolf [?] to estimate the correlation targeting matrix of the DCC model, and the second, which is original to this paper, is to calculate the log-likelihood of the DCC model using a series of rank-one updates to the cholesky decomposition of the quasi-correlation matrix at each time period, thus avoiding having to refactor this matrix at every time period.

The structure of this report is as follows: Section 2 will describe the Dynamic Conditional Correlation model, including the estimation problem; Section 3 will describe the linear shrinkage estimation for covariance matrices of Ledoit and Wolf.

2 Dynamic Conditional Correlation Model

In order to describe the Dynamic Conditional Correlation Model, it is necessary to build up some notation.

1. Let $r_{i,t}$ denote the return of asset i at time t , and r_t be the N -dimensional vector of all returns at time t .
2. Let $d_{i,t}^2 = V_{t-1}(r_{i,t}|\mathcal{F}_{t-1})$ be the conditional variance of asset i at time t , given information up to time $t-1$.
3. Let D_t be a diagonal matrix such that $D_{i,i,t} = d_{i,t}$.
4. Let H_t denote the conditional covariance matrix such that $H_{i,i,t} = \text{cov}_{t-1}(r_{i,t}, r_{j,t}|\mathcal{F}_{t-1})$.
5. Let $\epsilon_{i,t} = \frac{r_{i,t}}{d_{i,t}}$ be the standardised residuals of asset i at time t .
6. Let R_t be the conditional correlation matrix, whose i,j -th entry is given by the conditional correlation between asset i and asset j at time t .
7. Let σ_i^2 be the unconditional variance of the series $r_{i,t}$.
8. Let R be the unconditional correlation matrix of the system.

Assume we have N financial assets, and we observe for each of these assets a series of returns over a time period $1, \dots, T$.

The DCC model models financial assets jointly by first modelling their individual volatility structures and then modelling the correlation structure between them in a separate model.

Since the conditional correlation and conditional covariance are related by the equation

$$\rho_{i,j,t} = \frac{E_{t-1}((r_{i,t} - E_{t-1}(r_{i,t}))(r_{j,t} - E_{t-1}(r_{j,t})))}{\sqrt{V_{t-1}(r_{i,t})V_{t-1}(r_{j,t})}}, \quad (2.1)$$

these entries are therefore also given by

$$\frac{H_{i,j,t}}{\sqrt{H_{i,i,t}H_{j,j,t}}}. \quad (2.2)$$

Therefore, the conditional correlation matrix and conditional variance matrix are given by

$$R_t = D_t^{-1} H_t D_t^{-1} \quad D_t^2 = \text{diag}[H_t] \quad (2.3)$$

In the mean-reverting DCC model, the dynamics of the quasi-correlation matrix Q_t are governed by the process

$$Q_t = \Omega + \alpha \epsilon_{t-1} \epsilon'_{t-1} + \beta Q_{t-1}. \quad (2.4)$$

In the correlation targeting version of this model, the intercept matrix Ω is given

$$\Omega = (1 - \alpha - \beta) R \quad (2.5)$$

The matrix Q_t does not define a correlation matrix, as the diagonal elements are not necessarily 1. To convert the matrix Q_t into a correlation matrix we just divide the i, j -th entry by $\sqrt{Q_{i,i,t} Q_{j,j,t}}$, in other words

$$R_t = \text{diag}\{Q_t\}^{-\frac{1}{2}} Q_t \text{diag}\{Q_t\}^{-\frac{1}{2}} \quad (2.6)$$

This model is the *mean reverting* DCC model because we can see that over time as shocks to the correlation structure of the system decay, the system will tend towards its mean.

We will discuss in Section 4 the difficulty of estimating the parameters of this model when the number of assets considered, N , is large.

3 Shrinkage Estimation for Large Unconditional Correlation Matrices

As seen in Equation 2.3: correlation targeting, a key part of fitting the DCC model is estimating the unconditional correlation matrix R . We know from the cross-sectional literature that using the sample covariance matrix, in Equation 3.1, works poorly in large dimensions, since for large N , with $N \approx T$, the sample covariance matrix contains a lot of estimation error, since the correlation matrix has to estimate $\frac{N(N-1)}{2}$ parameters with T data points, so is prone to overfitting and provides a good fit in-sample but generalises poorly.

The *shrinkage* estimator of the covariance matrix aims to solve this problem of estimating the covariance matrix in high dimensions, essentially by first showing that the largest eigenvalues of the sample covariance matrix are larger than the true values, and the smallest eigenvalues are smaller than their true values. Therefore, if we *shrink* the sample eigenvalues towards their true value we should get a better estimate of the sample covariance matrix. Let Σ denote our true covariance matrix and $\hat{\Sigma}$ denote our sample covariance matrix. Further, let $\hat{\lambda}_i$ denote the i -th largest eigenvalue of $\hat{\Sigma}$.

$$\hat{\Sigma} = \frac{1}{T} \sum \epsilon_t \epsilon'_t \quad (3.1)$$

The *linear shrinkage* estimator of the covariance matrix, given by Ledoit and Wolf, aims to find the optimal linear combination

$$\Sigma^* = \rho_1 I + \rho_2 \hat{\Sigma} \quad (3.2)$$

such that the expected loss $\mathbb{E} \|\Sigma^* - \Sigma\|^2$ is minimal, where $\|A\|$ is the Frobenius norm of A , where for $A \in \mathbb{R}^{m \times n}$

$$\|A\| = \sqrt{\text{trace}(AA')}. \quad (3.3)$$

Ledoit and Wolf find an optimal, nonrandom solution for ρ_1, ρ_2 , but it requires prior knowledge about the true covariance matrix Σ . In particular, we need to be able to calculate the following values

$$\begin{aligned}\mu &= \langle \Sigma, I \rangle \\ \alpha^2 &= ||\Sigma - \mu I||^2 \\ \delta^2 &= \mathbb{E}||\hat{\Sigma} - \mu I||^2.\end{aligned}\tag{3.4}$$

Given these four scalar values, Ledoit and Wolf give a nonrandom optimal linear shrinkage estimator. Consider the problem we are trying to solve:

$$\begin{aligned}\min_{\rho_1, \rho_2} & \mathbb{E} [||\Sigma^* - \Sigma||^2] \\ \text{s.t. } & \Sigma^* = \rho_1 I + \rho_2 \hat{\Sigma}.\end{aligned}\tag{3.5}$$

Ledoit and Wolf show that this problem is solved by

$$\Sigma^* = \frac{\beta^2}{\delta^2} \mu I + \left(1 - \frac{\beta^2}{\delta^2}\right) \hat{\Sigma}\tag{3.6}$$

Obviously, however, we do not know the true covariance matrix, so we cannot directly compute these values and the optimal shrinkage intensity, *rho*. However, Ledoit and Wolf show the following constants can be combined to give a consistent estimator of the optimal shrinkage intensity.

$$\begin{aligned}m &= \langle \hat{\Sigma}, I \rangle \\ d^2 &= ||\hat{\Sigma} - mI||_n^2 \\ \bar{b}^2 &= \frac{1}{2} \sum_{k=1}^T ||\epsilon_t \epsilon'_t - \hat{\Sigma}||^2 \\ b^2 &= \min(\bar{b}^2, d^2)\end{aligned}\tag{3.7}$$

A consistent estimator for the optimal shrinking intensity $\hat{\rho}$ is then given by

$$\hat{\rho} = \frac{b^2}{d^2}\tag{3.8}$$

and then, our linear shrinkage estimate of the covariance matrix is

$$\bar{C} = \sum_{i=1}^N \left[\hat{\rho} \bar{\lambda} + (1 - \hat{\rho}) \hat{\lambda}_i \right] u_i u'_i\tag{3.9}$$

By a property of the Frobenius norm, the value μ is related to the eigenvalues of Σ , $\hat{\lambda}_i$ by the equation

$$\mu = \frac{1}{N} \sum_{i=1}^N \lambda_i = \mathbb{E} \left[\frac{1}{N} \sum_{i=1}^N \hat{\lambda}_i \right]\tag{3.10}$$

Ledoit and Wolf also show [?] that using a non-linear version of shrinkage estimation works better, but employing this technique is beyond the scope of this paper, and is an opportunity for further research. The main difference being that in nonlinear shrinkage the shrinkage parameter ρ is determined for each eigenvector individually rather than being forced to be the same for each eigenvector.

4 DCC Model Estimation

In Section

The problem of estimating the DCC model can be formulated as a maximum likelihood problem. In the DCC model where the conditional distribution of r_t

$$l = -\frac{1}{2} \sum_{t=1}^T n \log(2\pi) + 2 \log |D_t| + y_t' D_t^2 y_t - \epsilon_t' \epsilon_t + \log |R_t| + \epsilon_t' R_t^{-1} \epsilon_t \quad (4.1)$$

This log-likelihood can be split into two parts, the first part contains only the variance parameters from the GARCH specification in Equation 4.2, and the second contains only the DCC parameters, α and β . This log-likelihood can therefore be maximised in two stages, the first maximising with respect to the variance parameters $\omega_i, \alpha_i, \beta_i$, and the second stage maximising with respect to α, β . The first stage is done by fitting a univariate GARCH(1,1) model to each asset independently.

To model the quantities $d_{i,t}$, and from this extract the standardised residuals ϵ_t we apply the univariate GARCH(1,1) model to each asset separately to obtain the dynamics of its conditional variance.

$$H_{i,i,t} = \omega_i + \alpha_i r_{t-1}^2 + \beta_i H_{i,i,t-1}. \quad (4.2)$$

Then, the standardised residuals $\epsilon_{i,t}$ are given by

$$\epsilon_{i,t} = \frac{r_{i,t}}{d_{i,t}} \quad (4.3)$$

For the second stage, after fitting our univariate, independent GARCH(1,1) models, we extract the standardised residuals $\epsilon_{i,t}$ and use these to calculate the second part of the log-likelihood function

$$l_2 = -\frac{1}{2} \sum_{t=1}^T \log |R_t| + \epsilon_t' R_t^{-1} \epsilon_t. \quad (4.4)$$

To clarify how this relates to the specification of the mean-reverting DCC model in terms of α and β , let us remind ourselves of the dynamics of the quasi-correlation matrix Q_t , and how Q_t relates to the conditional correlation matrix R_t .

$$Q_t = \Omega + \alpha \epsilon_{t-1} \epsilon_{t-1}' + \beta Q_{t-1}, \quad (4.5)$$

$$\text{and } R_t = \text{diag}\{Q_t\}^{-\frac{1}{2}} Q_t \text{diag}\{Q_t\}^{-\frac{1}{2}}. \quad (4.6)$$

Therefore, in order to calculate the log-likelihood of our observed data, we need to, at every time interval t , invert the matrix R_t . Therefore, if T or N is large, directly estimating the log-likelihood in a naive way (by calculating the inverse of the matrix R_t for each $t \in \{1, \dots, T\}$) becomes computationally infeasible.

4.1 Cholesky Decomposition

The first thing to note is that we do not ever calculate the inverse of the matrix R_t . Rather we use the *Cholesky decomposition*, where, for a positive definite matrix A , $A = LL'$, where L is a lower triangular matrix.

$$L = \begin{bmatrix} L_{11} & 0 & 0 \\ L_{21} & L_{22} & 0 \\ L_{31} & L_{32} & L_{33} \end{bmatrix}. \quad (4.7)$$

This allows us to solve the system of linear equations $Ax = b$ (and thus $x = A^{-1}b$) using forward and back-substitution, without having to calculate A^{-1} . In particular, we want to, at each time interval t , calculate the quantity $R_t^{-1}\epsilon_t$. So, given the Cholesky decomposition of the matrix R_t , \tilde{L}_t we can compute $x_t = R_t^{-1}\epsilon_t$ and therefore $\epsilon_t'x_t = \epsilon_t'R_t^{-1}\epsilon_t$, the second part of l_2 . Therefore, if we can come up with a way of updating the cholesky decomposition of R_t at each time interval t , without having to refactor, then we can improve on the computational complexity of the log-likelihood calculation.

4.2 Updating the Cholesky Decomposition

Fortunately, there are algorithms for updating the cholesky decomposition of a matrix A when A is modified by

$$\tilde{A} = A + x'x \quad (4.8)$$

where x is a vector of length N . Such an update is called a rank one update, and the algorithm implemented for finding the Cholesky decomposition of \tilde{A} is given by Algorithm 1

Algorithm 1 Rank One Update to Cholesky decomposition

```

1: procedure RANKONEUP
2:   for  $i \in 1 : N$  do
3:      $r \leftarrow \sqrt{L_{ii}^2 + x_i^2}$ 
4:      $c \leftarrow r/L_{ii}$ 
5:      $s \leftarrow x_i/L_{ii}$ 
6:     for  $j \in i + 1 : N$  do
7:        $L_{ji} \leftarrow \frac{L_{ji} + sx_j}{c}$ 
8:        $x_j \leftarrow cx_j - sL_{ji}$ 
```

Noting that if we take x to be $\sqrt{\frac{\alpha}{\beta}}\epsilon_t$, this gives us an algorithm for finding the Cholesky decomposition of $Q_{t-1} + \frac{\alpha}{\beta}\epsilon_t\epsilon_t'$. Now note that if we multiply the resulting decomposition \tilde{L} by $\sqrt{\beta}$ we get the Cholesky decomposition of $\alpha\epsilon_t'\epsilon_t + \beta Q_{t-1}$. Therefore, this algorithm has given us an efficient way to update the Cholesky decomposition of the quasi-correlation matrix Q_t . However, the problem occurs when we add the full rank correlation target matrix Ω . There is no way of updating the Cholesky decomposition upon adding a full-rank matrix that is any more efficient than refactoring the matrix Q_t at every time interval t .

The innovation of this paper is to approximate each update by using the linear shrinkage estimator for the covariance matrix. In particular, rather than adding the full Ω matrix at each time step, we add the n largest eigenvalues. This allows us to estimate the true log-likelihood of some given DCC parameters, $\hat{\alpha}$ and $\hat{\beta}$, a process summarised in Algorithm 2, where the inputs are a matrix of returns, where r_{it} is the return of asset i at time t , and DCC parameter estimates α, β

If one were to update the matrix Q_t at each step, the algorithm would require $n^3/3$ FLOPs at each iteration to perform the Cholesky decomposition, and then a further

5 Monte Carlo Simulation

5.1 Setup of Simulations

To evaluate our method for approximating the log-likelihood of the DCC model, we use Monte Carlo simulation. In particular, we take as unconditional correlation matrix the Toeplitz matrix where the entries of

Algorithm 2 Estimating log-likelihood

```
1: procedure LOGLIK
2:   for  $i \in 1 : N$  do
3:     Fit GARCH(1,1) model to  $r_i$ 
4:      $\epsilon_{it} \leftarrow r_{it} / \hat{\sigma}_{it}$ 
5:      $\hat{\Sigma} \leftarrow \frac{1}{T} \sum_{t=1}^T \epsilon_t \epsilon_t'$ 
6:     Calculate eigenvalues, eigenvectors of  $\hat{\Sigma}$ 
7:      $\hat{\lambda}_i \leftarrow$   $i$ -th largest eigenvalue;  $u_i \leftarrow$  corresponding eigenvector
8:      $\bar{\lambda} \leftarrow \sum_{i=1}^N \hat{\lambda}_i$ 
9:     Calculate  $\hat{\rho}$  as in Equation 3.8
10:    for  $i \in 1 : N$  do
11:       $\lambda_i \leftarrow \hat{\rho} \bar{\lambda} + (1 - \hat{\rho}) \hat{\lambda}_i$ 
12:       $\bar{C} \leftarrow \frac{1}{T} \sum_{i=1}^T \lambda_i u_i u_i'$ 
13:       $Q_1 \leftarrow \bar{C}$ 
14:       $L \leftarrow \text{chol}(Q_1)$ 
15:       $l \leftarrow 0$ 
16:      for  $t \in 2 : T$  do
17:         $Q_t \leftarrow (1 - \alpha - \beta) \sum_{i=1}^n \lambda_i u_i u_i' + \alpha \epsilon_t \epsilon_t' + \beta Q_{t-1}$ 
18:         $L$  updated by Algorithm 1
19:         $l += \log |R_t| + \epsilon_t' R_t^{-1} \epsilon_t$  using  $L$ 
20:    return  $l$ 
```

the first row are given by the sequence $(1, \frac{n-1}{n}, \dots, \frac{1}{n})$ eg

$$L = \begin{bmatrix} 1 & \frac{2}{3} & \frac{1}{3} \\ \frac{1}{3} & 1 & \frac{2}{3} \\ \frac{2}{3} & \frac{1}{3} & 1 \end{bmatrix}. \quad (5.1)$$

We then simulate the DCC model, where the volatilities of each stock are governed by a GARCH(1,1) model with the parameterisation

$$\sigma_{it} = 0.01 + 0.05r_{t-1}^2 + 0.9\sigma_{i(t-1)}, \quad (5.2)$$

and the DCC parameters are $\alpha = 0.05$ and $\beta = 0.9$, as in Pakel et al [?] and Engle et al [?]. At each Monte Carlo simulation, we simulate a DCC model with the above dynamics, and then we look at the performance of our estimation method (using the estimated bias of the parameter estimates). We do this for $N \in \{5, 10, 50, 100, 200\}$ and for each of these, we choose a number of candidate values for n , ranging from 1 to N .

5.2 Results

Here we report the results from the Monte Carlo Simulations. As can be seen in Figures 5.2 and ??, increasing the number of rank one updates, n , to N gives the least biased estimates of α and β , with this bias being around $10\times$ higher for the n less than N . An interesting finding is that the bias of α increases as more updates are made. This is counterintuitive since as you add more eigenvectors you obtain a more accurate approximation of Ω , and so the estimated DCC parameters should be less biased. However, we see that in general, either using only a single rank-one update with the largest eigenvector or using N rank-one updates with all eigenvectors gives the least biased results.

N/n	5	10	50	100
1	0.031	0.027	0.024	0
5	-0.012	-0.035	-0.028	0
10	-	0.001	-0.049	0
50	-	-	-0.003	0
100	-	-	-	0

Figure 5.1: Bias of $\hat{\alpha}$ across simulations

N/n	5	10	50	100	200
1	-0.009	-0.001	-0.004	0	0
5	-0.018	0.055	0.048	0	0
10	-	-0.006	0.068	0	0
50	-	-	0.006	0.006	0
100	-	-	-	0	0

Figure 5.2: Bias of $\hat{\beta}$ across simulations

6 Conclusion

This thesis has proposed a method for efficiently estimating the parameters of a DCC model, using the linear shrinkage estimator of the covariance matrix as our correlation targeting matrix.




Figure 5.3: Estimated deviance for $N = 20$, $n = 20$

References

- [1] Engle, R. *Dynamic Conditional Correlation - a simple class of multivariate models*. Journal of Business and Economic Statistics, 2002.
- [2] Bollerslev, D *Dynamic Conditional Correlation - a simple class of multivariate models*. Journal of Business and Economic Statistics, 1990.
- [3] Engle, R. *Anticipating Correlations: a new paradigm for risk management*




Figure 5.4: Estimated deviance for $N = 20$, $n = 10$

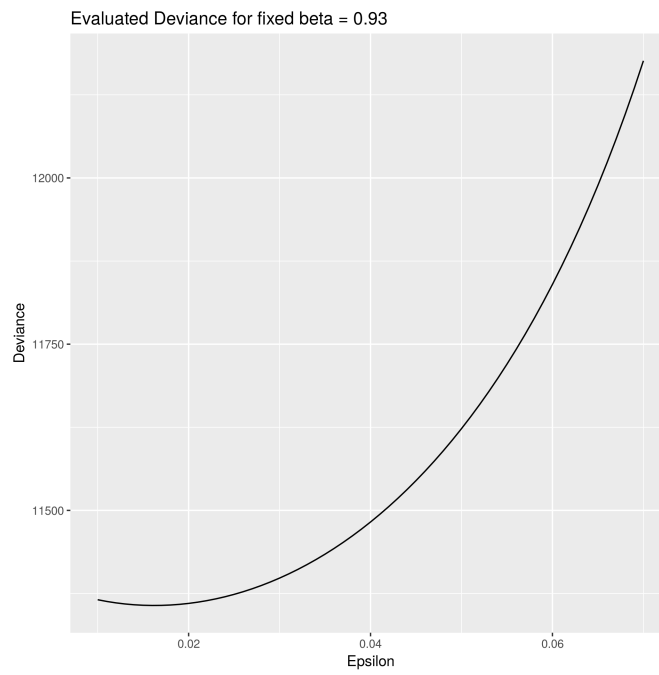


Figure 5.5: Estimated deviance for $N = 20$, $n = 1$